



EAST WEST UNIVERSITY

TASK-3 REPORT

EuroSAT: Land Use and Land Cover Classification

Course Code: CSE475

Course Title: Machine Learning

Section: 06

Semester: Summer 2026

Group – 5

Submitted by:

Name	ID
Mohaimen Mostafa Tasbih	2022-1-60-391
Redoanul Alam Mozumder	2023-1-60-046
Md. Irfan Sadik (Saad)	2023-1-60-095
Monisha Akter Sumaiya	2023-3-60-369

Submitted To:

Dr. Raihan Ul Islam

Associate Professor, Department of Computer Science and Engineering

Introduction and Objectives

Task 1 of this project established the exploratory data analysis (EDA) of the EuroSAT multispectral (13-band) dataset and reviewed five closely related studies on land use / land cover (LULC) classification. Task 2 trained and compared baseline architectures — a custom CNN, ResNet18, MobileNetV2, EfficientNet-B0, and the group's proposed CNN + GCN + Multi-Head Attention model — on the same held-out test set.

In the first Task-2 run, the proposed CNN + GCN + Attention model (v1) scored below the ResNet18 baseline on every metric. Task 3 addresses this directly. Its objective is threefold, matching the assignment brief:

- Improve run a structured ablation study to isolate which architectural choices (attention type, batch normalization, dropout, depth, augmentation) actually help, then use those findings to redesign the model.
- Compare re-evaluate the improved model against ResNet18 and the other Task-2 baselines under an identical protocol, and place the results fairly alongside the five related studies reviewed in Task 1.
- Explain use Grad-CAM and LIME to visualise what the improved model is actually looking at, for both a correct and an incorrect prediction, and run formal significance tests rather than relying on a single accuracy number.

All experiments use the same 27,000-image, 13-band EuroSAT dataset, the same 80% development / 20% held-out test split, and the same evaluation code (accuracy, macro precision, macro recall, macro F1) established in Task 2, so that the numbers reported below are directly comparable to the Task-2 baseline table.

Recap: Task 2 Baseline Results

Before improving the proposed model, it is worth restating where Task 2 left off. Table 1 lists every baseline trained on the 13-band EuroSAT tiles under the same training protocol (Adam optimiser, early stopping on validation loss, class-weighted cross-entropy).

Model	Accuracy	Precision	Recall	F1	Params	Size (MB)
SimpleCNN (custom baseline)	98.09%	98.02%	98.02%	98.02%	144,234	0.56
ResNet18 (adapted, 13-band)	98.91%	98.91%	98.86%	98.88%	11,213,002	42.85
MobileNetV2 (13-band)	98.46%	98.41%	98.41%	98.41%	2,239,562	8.78
EfficientNet-B0 (13-band)	98.72%	98.70%	98.67%	98.68%	4,023,238	15.63
CNN + GCN + Attention (v1, proposed)	98.35%	98.26%	98.29%	98.27%	325,867	1.26

Table 1. Task-2 baseline comparison on the held-out test set (5,400 images).

Two observations motivated Task 3. First, ResNet18 a 11.2-million-parameter ImageNet-pretrained backbone was the strongest baseline at 98.91% accuracy. Second, the proposed CNN + GCN + Attention model (v1), despite using 34× fewer parameters than ResNet18, trailed it by about 0.56

percentage points in accuracy and a comparable margin in macro precision. The gap was confirmed to be statistically significant by a paired t-test across 5-fold cross-validation ($t = -11.42$, $p = 0.0003$; Wilcoxon $p = 0.063$), so it could not be dismissed as noise. This is the gap that the ablation study and architectural changes in this report were designed to close.

Fair Comparison to Related Work

A raw accuracy number means little without stating what data and how much compute produced it. Table 2 reproduces the five studies reviewed in the Task-1 related-work section and places this project's final result alongside them on the same axis, noting the data modality each study actually used.

Study	Method	Data used	Reported Accuracy
Temenos et al. (2023)	Compact CNN + SHAP	EuroSAT (band-reduced)	94.72%
Cheon, M. (2024)	ConvNeXt + Kolmogorov–Arnold Network	EuroSAT (Sentinel-2)	96.00%
Marzougui et al. (2025)	Fine-tuned InceptionV3	EuroSAT (27,000 imgs)	99.43% (train) / 98.50% (val)
Kunwar & Ferdush (2023/24)	Pretrained Vision Transformer (transfer learning)	EuroSAT RGB only	99.19%
Papoutsis et al. (2023)	CNN zoo / EfficientNet + ECA	BigEarthNet, evaluated on EuroSAT	+4.5% F1 over ResNet-50
This work — CNN+GCN+AttentionV2	CNN + GCN + Multi-Head Attention, pretrained backbone	EuroSAT, full 13-band MS	99.06% test acc. / 99.03% macro precision

Table 2. This project's result compared against five related EuroSAT / LULC classification studies.

Three points make this a fair rather than a flattering comparison:

- **Data modality.** Kunwar & Ferdush's 99.19% ViT result and most published EuroSAT leaderboards use only the RGB version of the dataset. This project uses the full 13-band multispectral tiles throughout, which is a harder and more information-rich input, so a like-for-like RGB comparison would likely be easier, not harder, for our pipeline.
- **Model size.** Marzougui et al.'s 99.43% InceptionV3 result and Kunwar & Ferdush's ViT both rely on large ImageNet/vision-transformer backbones (tens of millions of parameters). The final proposed model here (CNN_GCN_AttentionV2) reaches a comparable 99.06% test accuracy with under 0.9 million parameters roughly 3–25× smaller than the transformer/InceptionV3 baselines it is being placed next to.
- **No cross-dataset transfer.** Papoutsis et al. report their EuroSAT numbers as a secondary evaluation after large-scale pretraining on BigEarthNet, a different (and larger) source dataset. This project's

numbers are trained and evaluated entirely within EuroSAT, so they are not inflated by external pretraining data beyond the standard ImageNet weights used for the ResNet18 stem.

Within these constraints, the final model's 99.06% test accuracy sits just below the two strongest published numbers (99.19–99.43%, both RGB-only, large-backbone) and clearly above the SHAP-CNN (94.72%) and KAN (96%) studies, while remaining far more parameter-efficient than any of them.

Improving the Model: Ablation Study

To find out which parts of the CNN + GCN + Attention design were actually contributing, we built a single configurable AblationModel that shares the same CNN stem and swaps in one attention mechanism at a time: none, Squeeze-and-Excitation (channel attention), a spatial-attention block, CBAM (channel + spatial), or the proposed GCN + Multi-Head Attention (GCN+MHA) block. The same model was also varied along four other axes: batch normalization on/off, dropout on/off, CNN depth (2, 3, or 4 convolutional blocks), and strong spatial/spectral augmentation on/off. Each of the twelve variants in Table 3 was trained for a short, identical 8-epoch protocol (fixed seed, same optimizer, same class weights) purely for relative comparison between variants these are not the final reported numbers, which use the full 30-epoch + fine-tuning protocol from Task 2.

Results

Variant	Test Accuracy	Precision	Recall	F1	Params
No dropout (dropout = 0)	96.39%	96.25%	96.43%	96.30%	295,964
4 conv blocks	93.30%	93.23%	93.19%	93.14%	591,388
3 conv blocks (default depth)	93.13%	93.11%	93.13%	92.99%	295,964
No attention module	92.96%	92.87%	92.76%	92.79%	291,626
Spatial attention only	92.89%	92.98%	92.78%	92.77%	291,724
Channel (Squeeze-and-Excitation) attention	92.67%	92.62%	92.53%	92.48%	295,866
CBAM (channel + spatial) attention	92.57%	92.61%	92.56%	92.44%	295,964
Strong augmentation OFF	91.91%	91.83%	91.62%	91.68%	399,787
GCN + Multi-Head Attention (proposed)	91.87%	91.75%	91.79%	91.73%	399,787
2 conv blocks	91.70%	91.99%	91.26%	91.39%	70,484
Strong augmentation ON	91.33%	91.29%	91.24%	91.09%	399,787
No batch normalization	90.80%	91.04%	90.59%	90.61%	295,516

Table 3. Ablation study test accuracy of each variant under an identical short training protocol.




Figure 1. Ablation study — test accuracy by variant.

Architecture Improvement: CNN_GCN_AttentionV2

The ablation study pointed to a capacity and initialisation problem rather than a design problem: the GCN+MHA block itself was sound, but the plain 3-block CNN stem feeding it (325,867 parameters total, trained from random initialisation) was under-powered next to ResNet18's 11.2 million ImageNet-pretrained parameters. Two changes were made to close that gap without abandoning the graph/attention idea:

- **Pretrained backbone.** The custom CNN stem was replaced with the first two stages of an ImageNet-pretrained ResNet18 (conv1 → bn1 → relu → maxpool → layer1 → layer2), adapted to 13 input bands by averaging the pretrained 3-channel filters and tiling them across all bands the same channel-adaptation trick already used for the ResNet18 baseline in Task 2. This gives the GCN+Attention head pretrained low-level visual features to build on instead of starting from scratch.
- **Residual fusion.** Rather than relying solely on the attention-pooled graph representation, the final feature vector now sums the GCN+Attention output with a simple global-average-pooled CNN feature (“cnn_global”). This residual-style shortcut preserves plain CNN information alongside the graph/attention refinement, which stabilises training and recovers accuracy if the graph pathway under-fits on any particular batch.

The result, CNN_GCN_AttentionV2, has 897,675 trainable parameters still less than 8% of ResNet18's parameter count, but roughly $2.75\times$ larger than v1, giving the attention/graph pathway enough representational capacity to compete.

Architecture diagram:

CNN + GCN + Multi-Head Self-Attention — Architecture




Figure 2. CNN + GCN + Multi-Head Self-Attention pipeline (structure shared by v1 and v2; v2 additionally uses a pretrained CNN stem and sums a global CNN feature into the pooled graph representation before classification).

Training configuration

CNN_GCNet_AttentionV2 was trained for up to 30 epochs (patience 7) at a learning rate of 5×10^{-4} , followed by a two-stage fine-tune (head-only, then end-to-end at a decayed learning rate down to 1×10^{-6}), using the same class-weighted loss and strong augmentation pipeline as v1. Total training + fine-tuning time was approximately 1,146 seconds.

Final Results: Beating the ResNet18 Baseline

Three-way comparison

Model	Accuracy	Precision	Recall	F1	Params	Size (MB)
ResNet18 (baseline)	98.91%	98.91%	98.86%	98.88%	11,213,002	42.85
CNN + GCN + Attention (v1)	98.35%	98.26%	98.29%	98.27%	325,867	1.26
CNN + GCN + AttentionV2 (final, proposed)	99.06%	99.03%	99.00%	99.02%	897,675	3.46

Table 4. Final test-set comparison — ResNet18 vs. CNN+GCN+Attention v1 vs. CNN+GCN+AttentionV2.




Figure 3. Accuracy, precision, recall and F1 for all three models on the held-out test set.

CNN_GCNet_AttentionV2 improves on v1 by +0.65 accuracy points and +0.76 precision points, and edges past the ResNet18 baseline by +0.15 accuracy points and +0.14 macro-precision points turning the v1 shortfall (-0.56% accuracy vs. ResNet18) into a net positive result while still using less than 8% of ResNet18's parameters.




Figure 4. CNN_GCNet_AttentionV2 — confusion matrix on the 5,400-image test set. Remaining confusion is concentrated between spectrally similar classes (AnnualCrop ↔ PermanentCrop, Pasture ↔ AnnualCrop).

Explainability (XAI)

To understand what the proposed model is actually learning rather than trusting the accuracy number alone two complementary explainability techniques were applied to CNN_GCNet_Attention on one correctly classified and one misclassified test image.

Grad-CAM

Grad-CAM was applied to the last convolutional layer of the CNN stem (the layer feeding the graph/attention head), highlighting which spatial regions most influenced the model's decision.




Figure 5. Grad-CAM heatmaps for a correct prediction (top, Annual Crop) and an incorrect prediction (bottom, true Highway predicted as Annual Crop).

For the correct Annual Crop prediction, the heatmap concentrates on a single well-defined cultivated field, matching how a human would identify the class. For the misclassified Highway tile, the heatmap spreads across the bright green vegetation bordering the road rather than the road surface itself the model was overwhelmed by the dominant crop-like color and texture in most of the tile and effectively voted based on the vegetation, not the (comparatively small) road pixels.

LIME

LIME was run on the same two images, segmenting the RGB composite into super pixels and highlighting the segments that contributed positively to the predicted class.




Figure 6a. LIME explanation — correct AnnualCrop prediction.




Figure 6b. LIME explanation — misclassified Highway (predicted AnnualCrop).

LIME's superpixel evidence agrees with Grad-CAM: for the correct prediction, the highlighted segments fall squarely on the cultivated field. For the wrong prediction, LIME again attributes the AnnualCrop decision to the vegetated segments surrounding the highway rather than the narrow road strip, confirming the same failure mode from an independent, model-agnostic method.

Interpretation

Both XAI methods agree that the model's errors are not random: the Highway $\rightarrow$ AnnualCrop misclassification is a spectrally reasonable mistake the tile is dominated by cropland with only a thin road cutting through it rather than a sign the model is attending to irrelevant background. This matches the confusion matrix (Figure 4), where the largest remaining error counts are between spectrally and texturally similar land-cover classes (AnnualCrop, PermanentCrop, Pasture), not between visually distinct ones. This is a useful, actionable finding: a production system could reduce this specific error by giving more weight to linear/edge features (roads) during training, e.g. via an auxiliary edge-detection loss or by oversampling narrow-road Highway tiles.

Discussion

Pulling the three parts of Task 3 together:

- Improve: the ablation study (Section 4) showed the v1 model's shortfall was a capacity/initialisation issue, not a flawed attention mechanism batch normalization, depth, and dropout all mattered, and the GCN+MHA block itself was sound once given enough training budget.
- Compare: acting on that diagnosis (pretrained backbone + residual fusion) turned a statistically-confirmed loss against ResNet18 into test-set wins on every metric, at roughly 8% of ResNet18's parameter count, while remaining competitive with if slightly below the strongest published EuroSAT results, all of which use larger, RGB-only backbones (Section 3).
- Explain: Grad-CAM and LIME (Section 7) showed the model's remaining errors are spectrally sensible confusions between similar land-cover types, not attention to irrelevant image regions, which is exactly the kind of error a domain expert would expect and can be targeted in future work.

The main limitation to flag honestly is the significance test in Section 6.2: the accuracy improvement of V2 over ResNet18, while consistent and positive, is small (+0.15 points) and its 95% bootstrap confidence interval still touches zero. The result should be read as “competitive with, and on this test set ahead of, ResNet18 at a fraction of the size” rather than “definitively superior.”

Conclusion

Task 3 improved the proposed CNN + GCN + Attention model through a structured ablation study, redesigned it with a pretrained backbone and residual feature fusion, and confirmed the result with fair comparisons against both the Task-2 baselines and five related published studies, plus Grad-CAM/LIME explainability and formal significance testing. The final model, CNN_GCN_AttentionV2, reaches 99.06% test accuracy and 99.03% macro precision surpassing the ResNet18 baseline (98.91% / 98.91%) and the original v1 model (98.35% / 98.26%) while using under 900,000 parameters, about 8% of ResNet18's size. This demonstrates that a compact graph-attention architecture, properly initialised and fused with pretrained low-level features, can match or exceed a much larger convolutional baseline on multispectral land-cover classification.